



Image Segmentation using U-Net on Oxford-IIIT Pet Dataset



- This project implements image segmentation using the U-Net architecture on the Oxford-IIIT Pet Dataset. Image segmentation is a computer vision task that involves labeling each pixel in an image with its corresponding class, allowing for detailed understanding of objects and their boundaries. In our case, we classify each pixel as either pet, background, or outline.



- The U-Net model, known for its U-shaped structure, consists of a contracting path to capture context and an expanding path to enable precise localization using skip connections. The Oxford-IIIT Pet Dataset contains over 7,000 images of 37 pet breeds with corresponding segmentation masks. The model processes these images to generate pixel-level predictions, enabling accurate pet segmentation with well-defined edges and minimal background noise.

Importing libraries

```
In [2]: import tensorflow as tf
import tensorflow.keras
import keras
import tensorflow_datasets as tfds
from tensorflow.keras import layers, Sequential, models
from keras.layers import Conv2D, MaxPooling2D, Dense, Flatten, Dropout
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
import pandas as pd
```

Loading the dataset

```
In [3]: dataset, info = tfds.load('oxford_iiit_pet', with_info=True)
```

Preprocessing

To prepare the Oxford-IIIT Pet Dataset for training, several preprocessing steps were applied. First, each image and its corresponding segmentation mask were resized to a fixed dimension of 128x128 pixels. The pixel values of the images were normalized to the range [0, 1], and the mask values were adjusted by subtracting 1 to ensure class labels start from zero. For data augmentation, a random horizontal flip was applied to increase the diversity of training samples and improve generalization. The dataset was then batched, shuffled with a buffer size of 1000, and prefetching was enabled to ensure optimal performance during training. These preprocessing steps help stabilize and speed up the training process while improving model robustness.

```
In [4]: len(dataset['train'])
```

```
Out[4]: 3680
```

```
In [5]: info
```

```

Out[5]: tfds.core.DatasetInfo(
    name='oxford_iiit_pet',
    full_name='oxford_iiit_pet/4.0.0',
    description="""
    The Oxford-IIIT pet dataset is a 37 category pet image dataset with rou
ghly 200
    images for each class. The images have large variations in scale, pose
and
    lighting. All images have an associated ground truth annotation of bree
d and
    species. Additionally, head bounding boxes are provided for the trainin
g split,
    allowing using this dataset for simple object detection tasks. In the t
est
    split, the bounding boxes are empty.
    """,
    homepage='http://www.robots.ox.ac.uk/~vgg/data/pets/',
    data_dir='/root/tensorflow_datasets/oxford_iiit_pet/4.0.0',
    file_format=tfrecord,
    download_size=773.52 MiB,
    dataset_size=773.68 MiB,
    features=FeaturesDict({
        'file_name': Text(shape=(), dtype=string),
        'head_bbox': BBoxFeature(shape=(4,), dtype=float32),
        'image': Image(shape=(None, None, 3), dtype=uint8),
        'label': ClassLabel(shape=(), dtype=int64, num_classes=37),
        'segmentation_mask': Image(shape=(None, None, 1), dtype=uint8),
        'species': ClassLabel(shape=(), dtype=int64, num_classes=2),
    }),
    supervised_keys=('image', 'label'),
    disable_shuffling=False,
    nondeterministic_order=False,
    splits={
        'test': <SplitInfo num_examples=3669, num_shards=4>,
        'train': <SplitInfo num_examples=3680, num_shards=4>,
    },
    citation="""@InProceedings{parkhil2a,
    author      = "Parkhi, O. M. and Vedaldi, A. and Zisserman, A. and J
awahar, C.~V.",
    title       = "Cats and Dogs",
    booktitle   = "IEEE Conference on Computer Vision and Pattern Recogn
ition",
    year        = "2012",
    }""",
)

```

```

In [6]: def normalize(input_image, input_mask):
    input_image = tf.cast(input_image, tf.float32)/255
    input_mask = input_mask-1
    return input_image, input_mask

```

```

In [7]: def load_train_images(sample):
    input_image= tf.image.resize(sample['image'],(128,128))
    input_mask = tf.image.resize(sample['segmentation_mask'],(128,128))
    if tf.random.uniform(())>0.5:

```

```

    input_image = tf.image.flip_left_right(input_image)
    input_mask = tf.image.flip_left_right(input_mask)
    input_image, input_mask = normalize(input_image, input_mask)
    return input_image , input_mask

```

```

In [8]: def load_test_images(sample):
        input_image= tf.image.resize(sample['image'],(128,128))
        input_mask = tf.image.resize(sample['segmentation_mask'],(128,128))
        input_image, input_mask = normalize(input_image, input_mask)
        return input_image , input_mask

```

```

In [9]: train_dataset = dataset['train'].map(load_train_images, num_parallel_calls=tf.data.experimental.AUTOTUNE)
        test_dataset = dataset['test'].map(load_test_images,num_parallel_calls=tf.data.experimental.AUTOTUNE)

```

```

In [10]: BATCH_SIZE=64
        BUFFER_SIZE=1000
        train_dataset = train_dataset.cache().shuffle(BUFFER_SIZE).batch(BATCH_SIZE)
        train_dataset = train_dataset.prefetch(buffer_size=tf.data.experimental.AUTOTUNE)
        test_dataset = test_dataset.batch(BATCH_SIZE)

```

EDA

```

In [11]: def display_sample(image_list):
        plt.figure(figsize=(8,8))
        title = ['Input Image', 'True Mask' , 'Predicted Mask']
        for i in range(len(image_list)):
            plt.subplot(1,len(image_list),i+1)
            plt.title(title[i])
            plt.imshow(tf.keras.preprocessing.image.array_to_img(image_list[i]))
            plt.axis('off')
        plt.show()

```

```

In [12]: for images, masks in train_dataset.take(3):
        sample_image, sample_mask = images[0], masks[0]
        display_sample([sample_image, sample_mask])

```


Input Image



True Mask



Input Image



True Mask



Input Image



True Mask



In [12]:

Define U-Net Model

```
In [13]: def double_conv_block(x, n_filters):
    x= Conv2D(n_filters, 3, padding='same', activation='relu', kernel_initializer='he_normal')(x)
    x= Conv2D(n_filters, 3, padding='same', activation='relu', kernel_initializer='he_normal')(x)
    return x

def downsample_block(x, n_filters):
    f = double_conv_block(x, n_filters)
    p = MaxPooling2D((2,2))(f)
    p = Dropout(0.3)(p)
    return f, p

def upsample_block(x, conv_features, n_filters):
    x= layers.Conv2DTranspose(n_filters, 3, 2, padding='same')(x)
    x= layers.Concatenate()([x, conv_features])
    x= Dropout(0.3)(x)
    x = double_conv_block(x, n_filters)
    return x
```

```
In [14]: def build_unet_model(output_channels):
    inputs= layers.Input(shape=(128,128,3))
    f1, p1 = downsample_block(inputs, 64)
    f2, p2 = downsample_block(p1, 128)
    f3, p3 = downsample_block(p2, 256)
    f4, p4 = downsample_block(p3, 512)
    bottleneck = double_conv_block(p4, 1024)
    u6 = upsample_block(bottleneck, f4, 512)
    u7 = upsample_block(u6, f3, 256)
    u8 = upsample_block(u7, f2, 128)
    u9 = upsample_block(u8, f1, 64)
    outputs = layers.Conv2D(output_channels, 1, padding='same', activation='sigmoid')(u9)
    unet_model = tf.keras.Model(inputs, outputs, name='U-Net')
    return unet_model
```

```
In [15]: output_channels=3
model= build_unet_model(output_channels)
model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])
```

```
In [16]: model.summary()
```

Model: "U-Net"

Layer (type)	Output Shape	Param #	Connected to
input_layer (InputLayer)	(None, 128, 128, 3)	0	-
conv2d (Conv2D)	(None, 128, 128, 64)	1,792	input_layer[0]
conv2d_1 (Conv2D)	(None, 128, 128, 64)	36,928	conv2d[0][0]
max_pooling2d (MaxPooling2D)	(None, 64, 64, 64)	0	conv2d_1[0][0]
dropout (Dropout)	(None, 64, 64, 64)	0	max_pooling2d
conv2d_2 (Conv2D)	(None, 64, 64, 128)	73,856	dropout[0][0]
conv2d_3 (Conv2D)	(None, 64, 64, 128)	147,584	conv2d_2[0][0]
max_pooling2d_1 (MaxPooling2D)	(None, 32, 32, 128)	0	conv2d_3[0][0]
dropout_1 (Dropout)	(None, 32, 32, 128)	0	max_pooling2d
conv2d_4 (Conv2D)	(None, 32, 32, 256)	295,168	dropout_1[0][0]
conv2d_5 (Conv2D)	(None, 32, 32, 256)	590,080	conv2d_4[0][0]
max_pooling2d_2 (MaxPooling2D)	(None, 16, 16, 256)	0	conv2d_5[0][0]
dropout_2 (Dropout)	(None, 16, 16, 256)	0	max_pooling2d
conv2d_6 (Conv2D)	(None, 16, 16, 512)	1,180,160	dropout_2[0][0]
conv2d_7 (Conv2D)	(None, 16, 16, 512)	2,359,808	conv2d_6[0][0]
max_pooling2d_3 (MaxPooling2D)	(None, 8, 8, 512)	0	conv2d_7[0][0]
dropout_3 (Dropout)	(None, 8, 8, 512)	0	max_pooling2d
conv2d_8 (Conv2D)	(None, 8, 8, 1024)	4,719,616	dropout_3[0][0]
conv2d_9 (Conv2D)	(None, 8, 8, 1024)	9,438,208	conv2d_8[0][0]
conv2d_transpose	(None, 16, 16, 1024)	4,719,104	conv2d_9[0][0]

(Conv2DTranspose)	512)		
concatenate (Concatenate)	(None, 16, 16, 1024)	0	conv2d_transp conv2d_7[0][0
dropout_4 (Dropout)	(None, 16, 16, 1024)	0	concatenate[0
conv2d_10 (Conv2D)	(None, 16, 16, 512)	4,719,104	dropout_4[0][
conv2d_11 (Conv2D)	(None, 16, 16, 512)	2,359,808	conv2d_10[0][
conv2d_transpose_1 (Conv2DTranspose)	(None, 32, 32, 256)	1,179,904	conv2d_11[0][
concatenate_1 (Concatenate)	(None, 32, 32, 512)	0	conv2d_transp conv2d_5[0][0
dropout_5 (Dropout)	(None, 32, 32, 512)	0	concatenate_1
conv2d_12 (Conv2D)	(None, 32, 32, 256)	1,179,904	dropout_5[0][
conv2d_13 (Conv2D)	(None, 32, 32, 256)	590,080	conv2d_12[0][
conv2d_transpose_2 (Conv2DTranspose)	(None, 64, 64, 128)	295,040	conv2d_13[0][
concatenate_2 (Concatenate)	(None, 64, 64, 256)	0	conv2d_transp conv2d_3[0][0
dropout_6 (Dropout)	(None, 64, 64, 256)	0	concatenate_2
conv2d_14 (Conv2D)	(None, 64, 64, 128)	295,040	dropout_6[0][
conv2d_15 (Conv2D)	(None, 64, 64, 128)	147,584	conv2d_14[0][
conv2d_transpose_3 (Conv2DTranspose)	(None, 128, 128, 64)	73,792	conv2d_15[0][
concatenate_3 (Concatenate)	(None, 128, 128, 128)	0	conv2d_transp conv2d_1[0][0
dropout_7 (Dropout)	(None, 128, 128, 128)	0	concatenate_3
conv2d_16 (Conv2D)	(None, 128, 128, 64)	73,792	dropout_7[0][
conv2d_17 (Conv2D)	(None, 128, 128, 64)	36,928	conv2d_16[0][
conv2d_18 (Conv2D)	(None, 128, 128,	195	conv2d_17[0][

	3)		
--	----	--	--


Total params: 34,513,475 (131.66 MB)


Trainable params: 34,513,475 (131.66 MB)


Non-trainable params: 0 (0.00 B)


```
In [17]: EPOCHS=20
steps_per_epoch = info.splits['train'].num_examples // BATCH_SIZE
validation_steps = info.splits['test'].num_examples // BATCH_SIZE


history = model.fit(train_dataset, epochs=EPOCHS, steps_per_epoch=steps_per_
```


Epoch 1/20
57/57  **163s** 1s/step - accuracy: 0.5696 - loss: 1.0117 - val_accuracy: 0.5726 - val_loss: 0.8488


Epoch 2/20
57/57  **123s** 1s/step - accuracy: 0.5816 - loss: 0.8483 - val_accuracy: 0.5971 - val_loss: 0.8044


Epoch 3/20
57/57  **62s** 1s/step - accuracy: 0.6104 - loss: 0.7942 - val_accuracy: 0.6256 - val_loss: 0.7931


Epoch 4/20
57/57  **82s** 1s/step - accuracy: 0.6329 - loss: 0.7790 - val_accuracy: 0.6856 - val_loss: 0.8682


Epoch 5/20
57/57  **62s** 1s/step - accuracy: 0.6423 - loss: 0.7791 - val_accuracy: 0.6271 - val_loss: 0.7735


Epoch 6/20
57/57  **63s** 1s/step - accuracy: 0.6336 - loss: 0.7783 - val_accuracy: 0.6809 - val_loss: 0.7202


Epoch 7/20
57/57  **63s** 1s/step - accuracy: 0.7128 - loss: 0.6705 - val_accuracy: 0.7364 - val_loss: 0.6243


Epoch 8/20
57/57  **63s** 1s/step - accuracy: 0.7458 - loss: 0.6089 - val_accuracy: 0.7725 - val_loss: 0.5510


Epoch 9/20
57/57  **63s** 1s/step - accuracy: 0.7725 - loss: 0.5518 - val_accuracy: 0.7851 - val_loss: 0.5161


Epoch 10/20
57/57  **63s** 1s/step - accuracy: 0.7878 - loss: 0.5163 - val_accuracy: 0.7965 - val_loss: 0.5034


Epoch 11/20
57/57  **63s** 1s/step - accuracy: 0.7987 - loss: 0.4934 - val_accuracy: 0.8134 - val_loss: 0.4621


Epoch 12/20
57/57  **82s** 1s/step - accuracy: 0.8150 - loss: 0.4521 - val_accuracy: 0.8228 - val_loss: 0.4412


Epoch 13/20
57/57  **82s** 1s/step - accuracy: 0.8246 - loss: 0.4323 - val_accuracy: 0.8193 - val_loss: 0.4484


Epoch 14/20
57/57  **62s** 1s/step - accuracy: 0.8278 - loss: 0.4237 - val_accuracy: 0.8407 - val_loss: 0.3947

Epoch 15/20
57/57  **82s** 1s/step - accuracy: 0.8404 - loss: 0.3906 - val_accuracy: 0.8447 - val_loss: 0.3830

Epoch 16/20
57/57  **63s** 1s/step - accuracy: 0.8474 - loss: 0.3716 - val_accuracy: 0.8459 - val_loss: 0.3779

Epoch 17/20
57/57  **63s** 1s/step - accuracy: 0.8486 - loss: 0.3705 - val_accuracy: 0.8508 - val_loss: 0.3768

Epoch 18/20
57/57  **82s** 1s/step - accuracy: 0.8521 - loss: 0.3595 - val_accuracy: 0.8482 - val_loss: 0.3790

Epoch 19/20
57/57  **82s** 1s/step - accuracy: 0.8613 - loss: 0.3372 - val_accuracy: 0.8613 - val_loss: 0.3372

al_accuracy: 0.8537 - val_loss: 0.3676

Epoch 20/20

57/57 ————— 62s 1s/step - accuracy: 0.8640 - loss: 0.3296 - v


al_accuracy: 0.8430 - val_loss: 0.4041

In []:

Visualize the results


```
In [20]: ## plot train and val accuracy
plt.figure(figsize=(12,4))
plt.subplot(1,2,1)
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend(['train', 'val'], loc='upper left')
```

Out[20]: <matplotlib.legend.Legend at 0x7d9b37abacd0>



```
In [22]: plt.figure(figsize=(12,4))
plt.subplot(1,2,1)
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend(['train', 'val'], loc='upper left')
```

```
plt.tight_layout()
plt.show()
```



```
In [31]: def create_mask(pred_msk):
    pred_mask = tf.argmax(pred_msk, axis=-1)
    pred_mask = pred_mask[..., tf.newaxis]
    return pred_mask[0]

    def show_predictions(dataset=None, num=1):
        if dataset:
            for image, mask in dataset.take(num):
                pred_mask = model.predict(image)
                display = [image[0], mask[0], create_mask(pred_mask)]
                display_sample(display)
```

```
In [32]: for image, mask in test_dataset.take(1):
    pred_mask = model.predict(image)


    pred_mask = tf.argmax(pred_mask, axis=-1)
    pred_mask = pred_mask[..., tf.newaxis]
    print(pred_mask[0].shape)
```

```
2/2 ————— 0s 175ms/step
(128, 128, 1)
```

```
In [33]: show_predictions(test_dataset, 10)
```

```
2/2 ————— 0s 182ms/step
```

Input Image



True Mask



Predicted Mask



2/2

0s 166ms/step

Input Image



True Mask




Predicted Mask



2/2

0s 167ms/step

Input Image



True Mask




Predicted Mask



2/2

0s 165ms/step

Input Image



True Mask




Predicted Mask



2/2

0s 159ms/step

Input Image



True Mask




Predicted Mask



2/2  0s 158ms/step

Input Image



True Mask




Predicted Mask



2/2  0s 161ms/step

Input Image



True Mask



Predicted Mask



2/2  0s 161ms/step

Input Image



True Mask



Predicted Mask



2/2  0s 167ms/step



2/2 ————— 0s 160ms/step



```
In [38]: model.save('history.h5')
```

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

Model Evaluation

```
In [36]: loss, accuracy = model.evaluate(test_dataset)
print(f"Test Loss: {loss}")
print(f"Test Accuracy: {accuracy}")
```

58/58 ————— 26s 438ms/step - accuracy: 0.8434 - loss: 0.4016
 Test Loss: 0.40391096472740173
 Test Accuracy: 0.8430185914039612

```
In [37]: from tensorflow.keras.metrics import MeanIoU

# Pick the number of classes (in Oxford-IIIT Pet, usually 3: pet, background)
num_classes = 3
miou = MeanIoU(num_classes=num_classes)

for images, masks in test_dataset:
    preds = model.predict(images)
    preds = tf.argmax(preds, axis=-1)
    preds = preds[..., tf.newaxis]
    miou.update_state(masks, preds)
```

```
print("Mean IoU:", miou.result().numpy())
```


2/2	0s	177ms/step
2/2	0s	156ms/step
2/2	0s	157ms/step
2/2	0s	157ms/step
2/2	0s	162ms/step
2/2	0s	158ms/step
2/2	0s	163ms/step
2/2	0s	176ms/step
2/2	0s	169ms/step
2/2	0s	157ms/step
2/2	0s	156ms/step
2/2	0s	158ms/step
2/2	0s	158ms/step
2/2	0s	158ms/step
2/2	0s	158ms/step
2/2	0s	158ms/step
2/2	0s	161ms/step
2/2	0s	161ms/step
2/2	0s	158ms/step
2/2	0s	157ms/step
2/2	0s	159ms/step
2/2	0s	158ms/step
2/2	0s	158ms/step
2/2	0s	157ms/step
2/2	0s	157ms/step
2/2	0s	158ms/step
2/2	0s	166ms/step
2/2	0s	166ms/step
2/2	0s	159ms/step
2/2	0s	158ms/step
2/2	0s	157ms/step
2/2	0s	157ms/step
2/2	0s	162ms/step
2/2	0s	159ms/step
2/2	0s	158ms/step
2/2	0s	158ms/step
2/2	0s	157ms/step
2/2	0s	159ms/step
2/2	0s	159ms/step
2/2	0s	158ms/step
2/2	0s	158ms/step
2/2	0s	160ms/step
2/2	0s	158ms/step
2/2	0s	158ms/step
2/2	0s	163ms/step
2/2	0s	159ms/step
2/2	0s	162ms/step
2/2	0s	181ms/step
2/2	0s	155ms/step
2/2	0s	159ms/step
2/2	0s	159ms/step
2/2	0s	159ms/step
2/2	0s	159ms/step
2/2	0s	160ms/step
2/2	0s	158ms/step
2/2	0s	156ms/step
2/2	0s	160ms/step

2/2  0s 158ms/step

1/1  1s 959ms/step

Mean IoU: 0.6517233



Model Performance Summary

The U-Net-based image segmentation model trained on the Oxford-IIIT Pet Dataset demonstrates strong performance in pixel-wise classification:

Test Loss: 0.4039

Test Accuracy: 84.30%

Mean Intersection over Union (IoU): 0.65

These metrics indicate that the model has effectively learned to distinguish between the foreground (pet), background, and outlines, achieving a reliable segmentation quality across the validation samples.



Conclusion

The U-Net architecture, known for its encoder-decoder structure and skip connections, proved to be highly effective for the task of semantic segmentation. With a test accuracy of over 84% and a mean IoU of 0.65, the model demonstrates a good balance between precision and recall in segmenting pet images. This validates the robustness of the chosen preprocessing pipeline and training strategy, making the model suitable for real-world image segmentation tasks where spatial accuracy is critical.

In []:

This notebook was converted with convert.ploomber.io